

# Generative Adversarial Networks (GANs)

Adapted from material by Goodfellow, Binglin, Shashank, Bhargav

# GANs

- **Generative**

- Learn a generative model

- **Adversarial**

- Trained in an adversarial setting

- **Networks**

- Use Deep Neural Networks

# Why Generative Models?

- **Discriminative models:**
  - Given an image  $\mathbf{X}$ , predict a label  $\mathbf{Y}$
  - Estimates  $\mathbf{P}(\mathbf{Y}|\mathbf{X})$
- **Discriminative models have several key limitations**
  - Can't model  $\mathbf{P}(\mathbf{X})$ , i.e. the probability of seeing a certain image
  - Thus, can't sample from  $\mathbf{P}(\mathbf{X})$ , i.e. **can't generate new images**
- **Generative models try to address these:**
  - model  $\mathbf{P}(\mathbf{X})$
  - generate new data (e.g. images)

# Magic of GANs...

Which one is Computer generated?



# Magic of GANs...

User edits



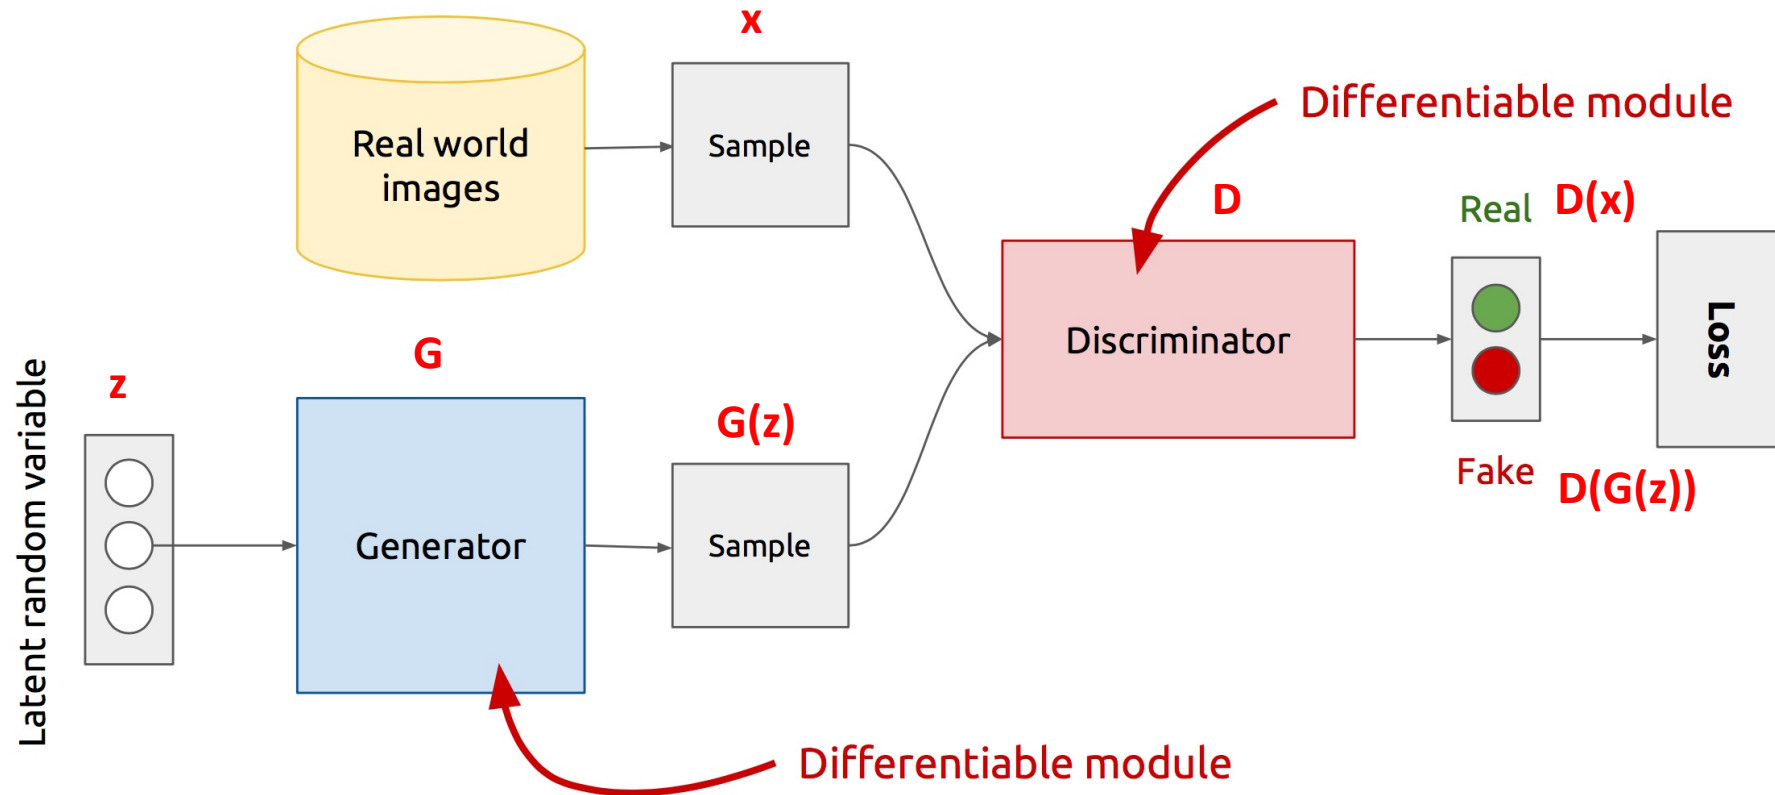
Generated images



# Adversarial Training

- Generator: generate fake samples, tries to fool the Discriminator
- Discriminator: tries to distinguish between real and fake samples
- Train them against each other
- Repeat this and we get better Generator and Discriminator

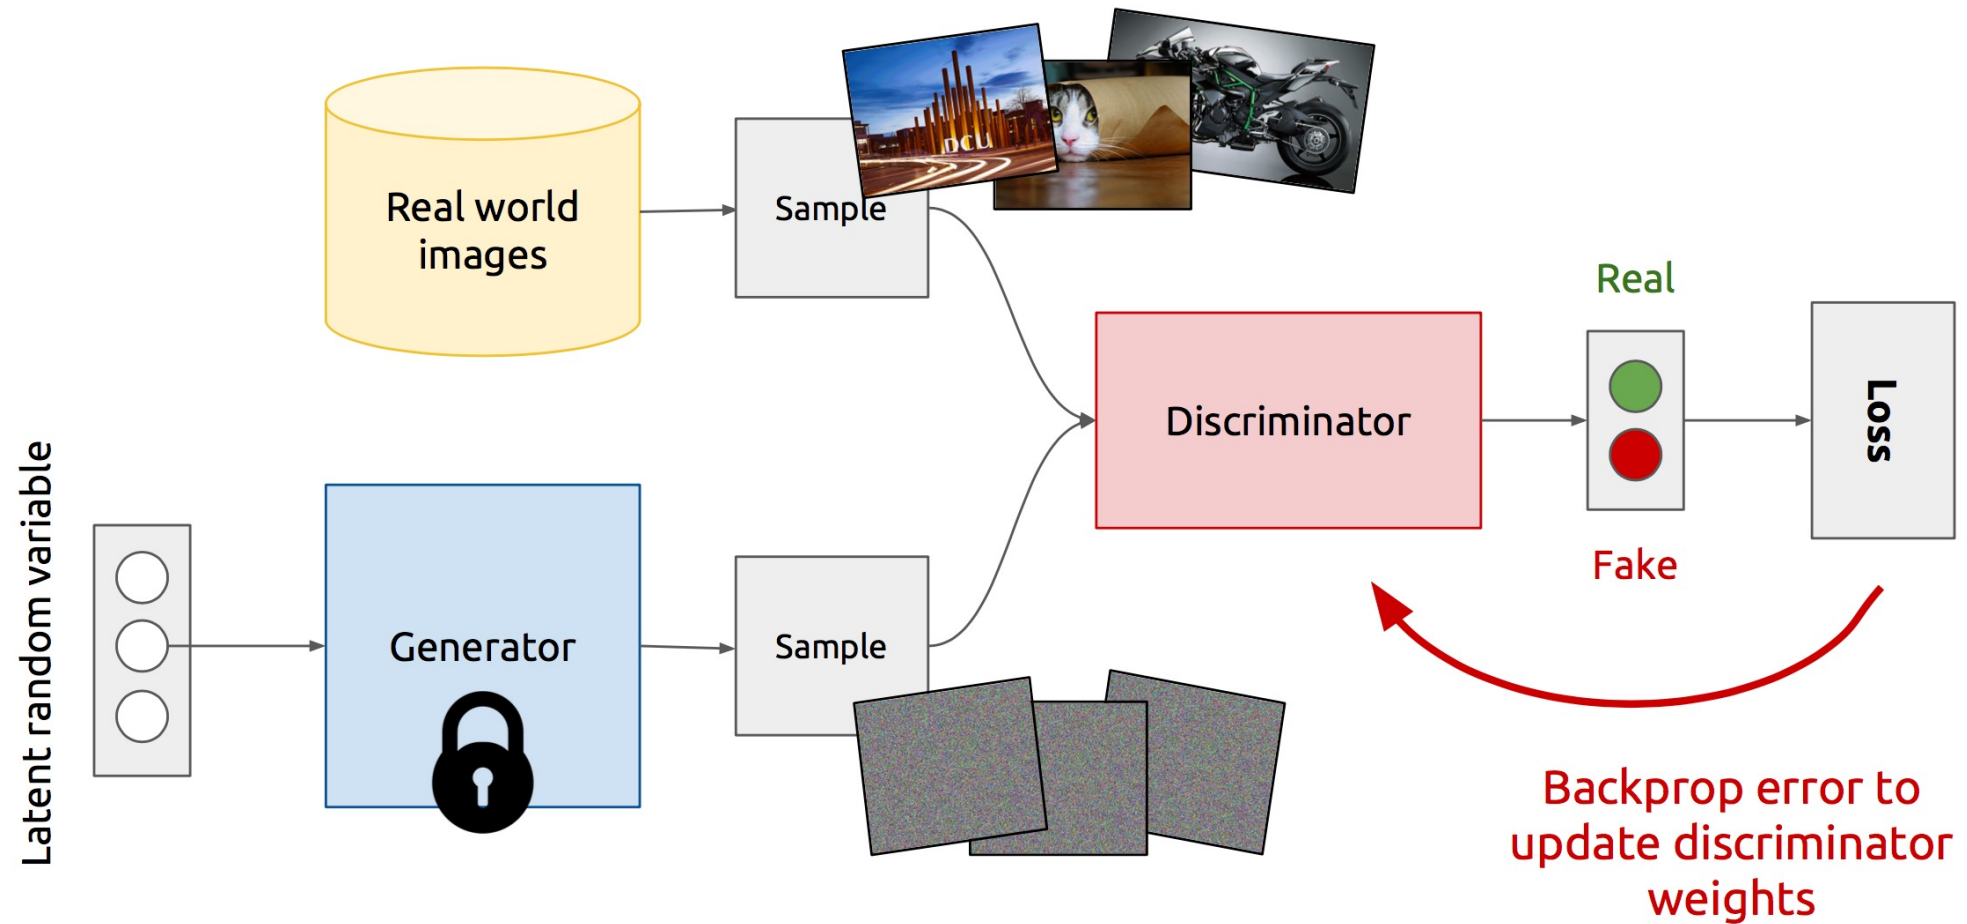
# GAN's Architecture



- **Z** is some random vector (Gaussian/Uniform).
- **Z** can be thought as the latent representation of the image.

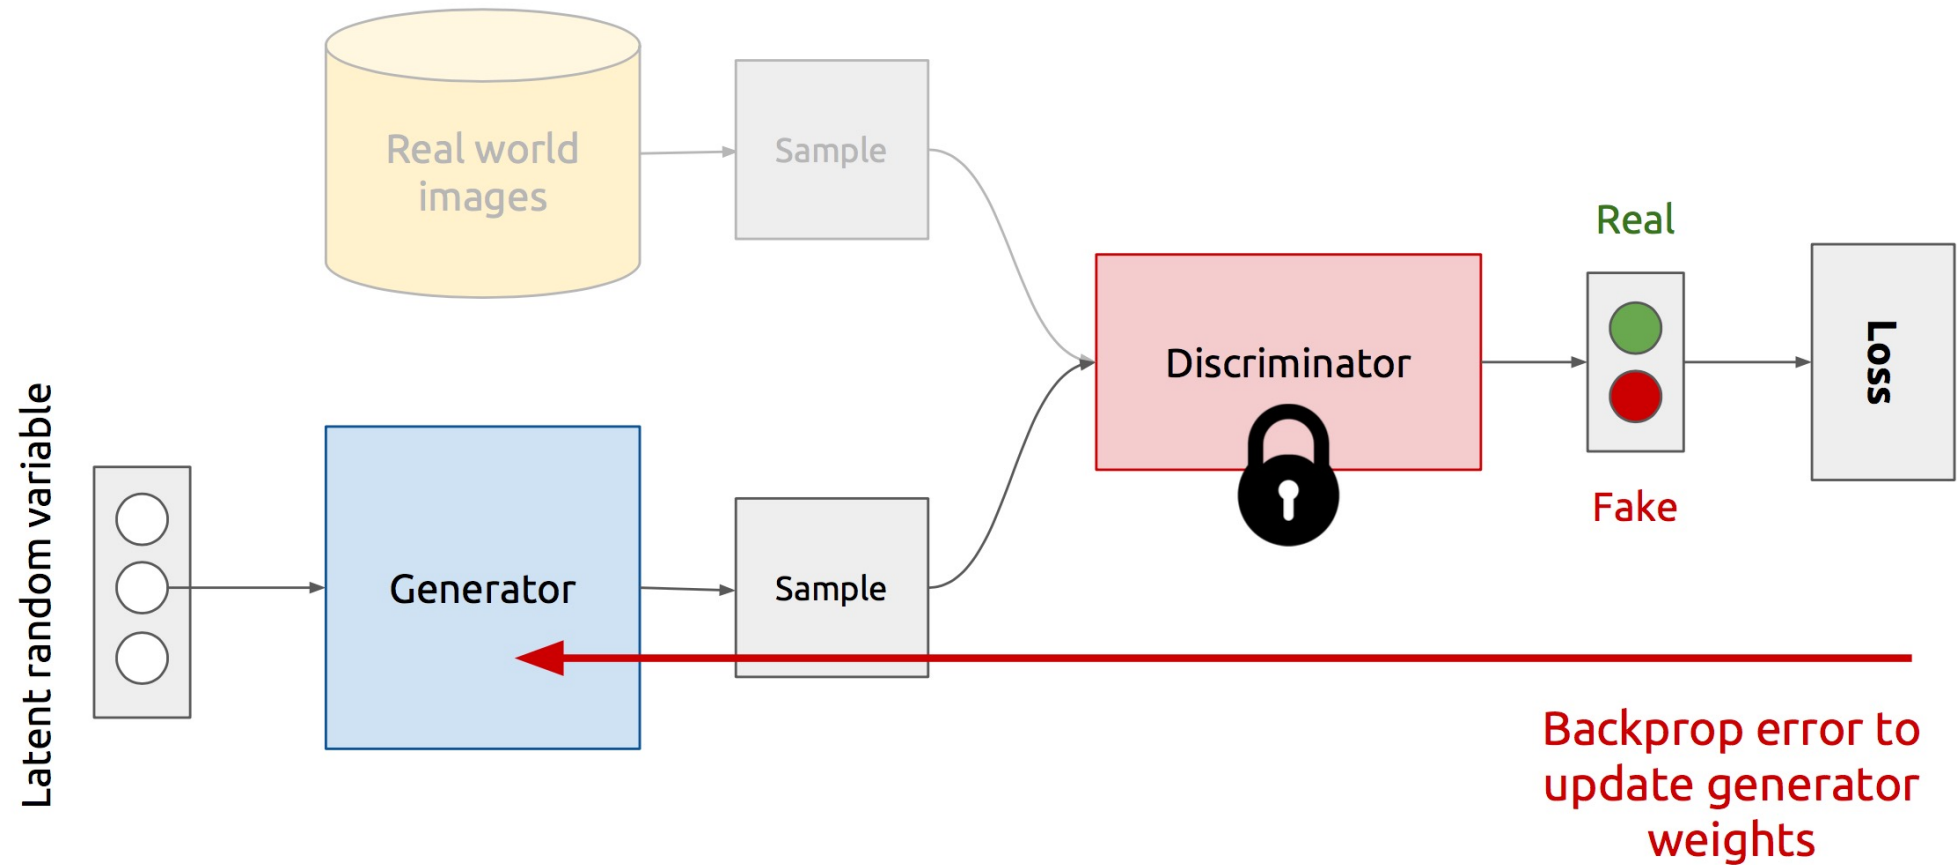


# Training Discriminator





# Training Generator



# GAN's formulation

$$\min_G \max_D V(D, G)$$

- It is formulated as a **minimax game**, where:
  - The Discriminator is trying to maximize its reward  $V(D, G)$
  - The Generator is trying to minimize Discriminator's reward (or maximize its loss)

$$V(D, G) = \mathbb{E}_{x \sim p(x)} [\log D(x)] + \mathbb{E}_{z \sim q(z)} [\log(1 - D(G(z)))]$$

- The Nash equilibrium of this particular game is achieved at:
  - $P_{data}(x) = P_{gen}(x) \quad \forall x$
  - $D(x) = \frac{1}{2} \quad \forall x$

---

**Algorithm 1** Minibatch stochastic gradient descent training of generative adversarial nets. The number of steps to apply to the discriminator,  $k$ , is a hyperparameter. We used  $k = 1$ , the least expensive option, in our experiments.

---

**for** number of training iterations **do**

**for**  $k$  steps **do**

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Sample minibatch of  $m$  examples  $\{x^{(1)}, \dots, x^{(m)}\}$  from data generating distribution  $p_{\text{data}}(x)$ .
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[ \log D(x^{(i)}) + \log (1 - D(G(z^{(i)}))) \right].$$

**end for**

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Update the generator by descending its stochastic gradient:

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log (1 - D(G(z^{(i)}))).$$

**end for**

The gradient-based updates can use any standard gradient-based learning rule. We used momentum in our experiments.

---


**Discriminator  
updates**

**Generator  
updates**

# Faces



# CIFAR



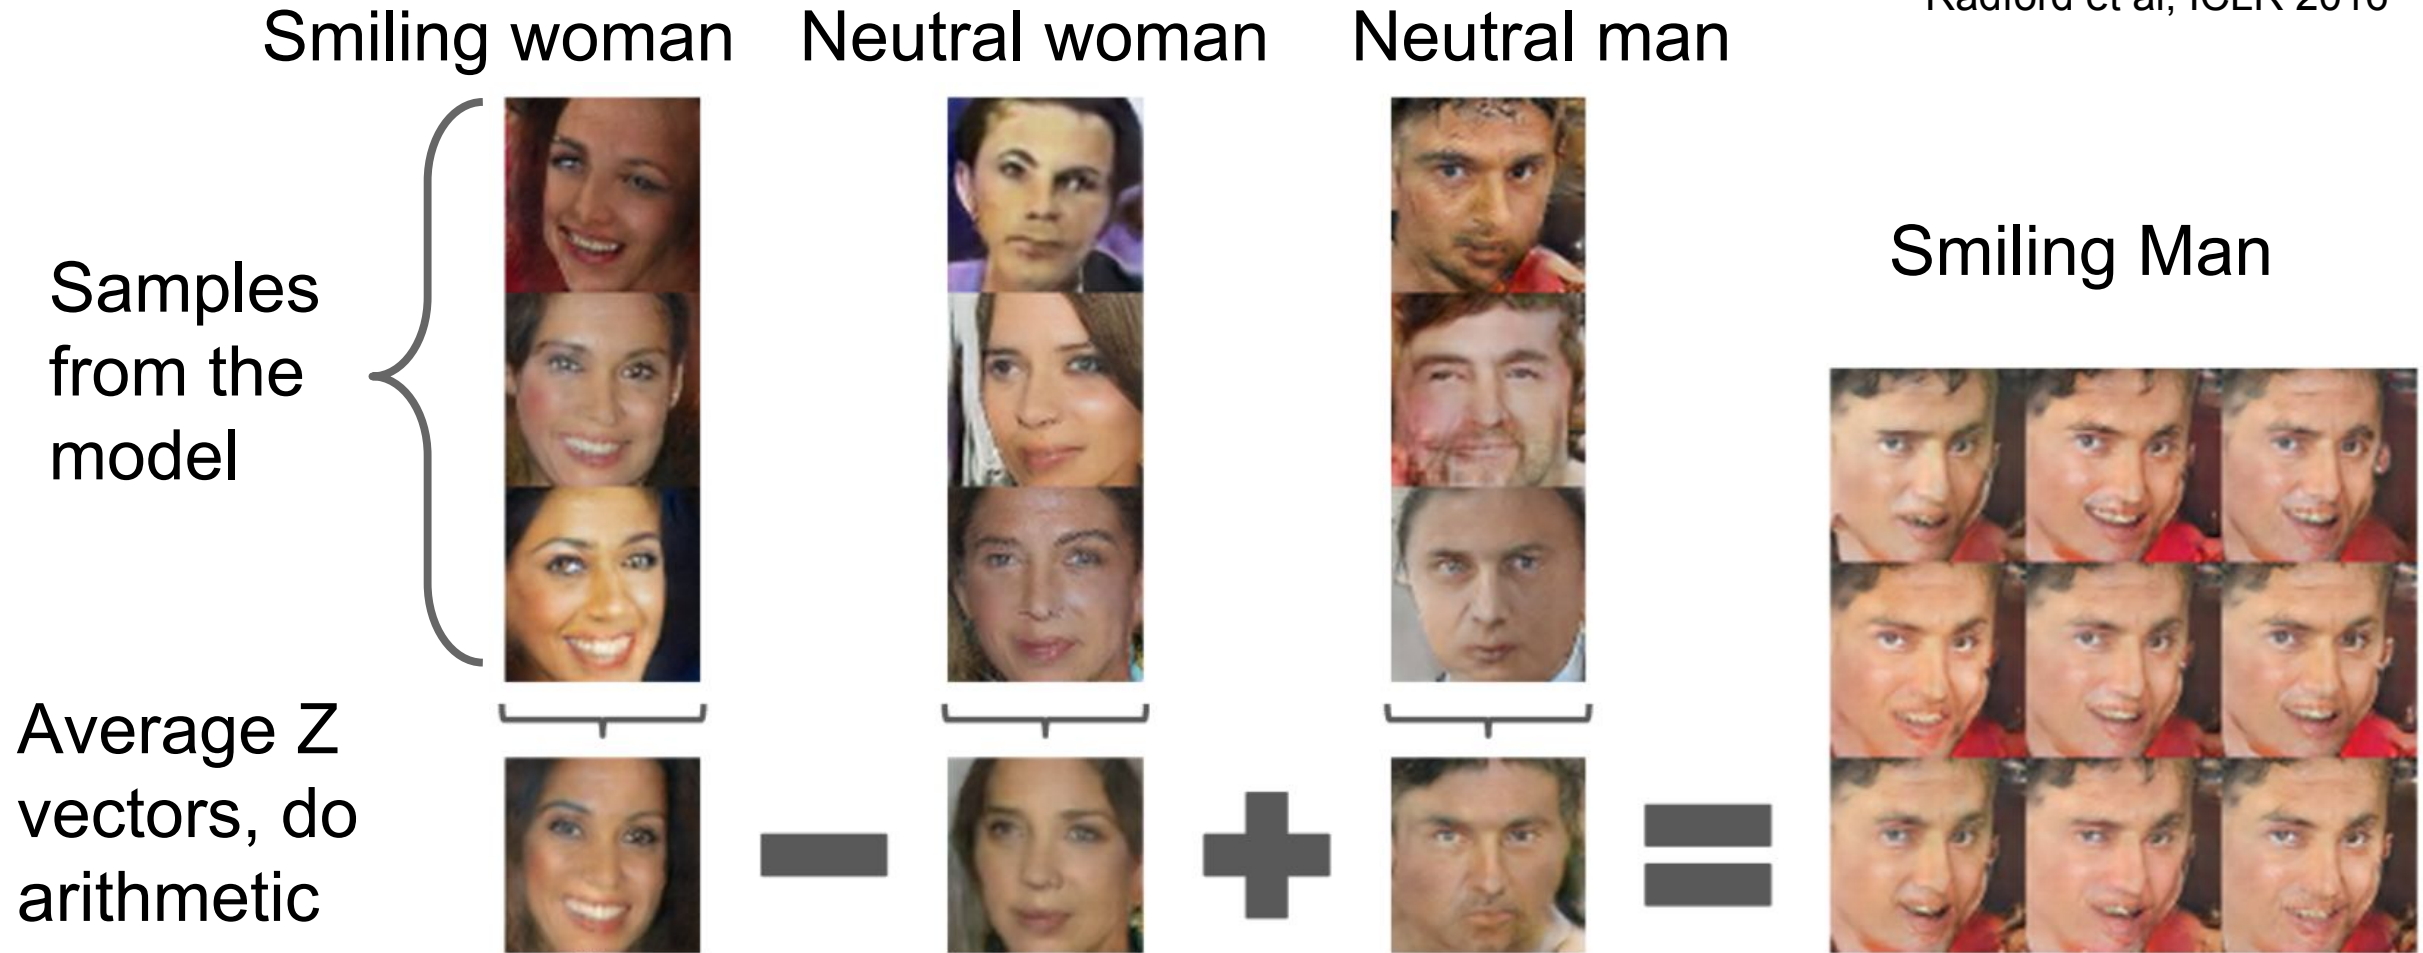


# DCGAN: Bedroom images



# Generative Adversarial Nets: Interpretable Vector Math

Radford et al, ICLR 2016





# Generative Adversarial Nets: Interpretable Vector Math

Glasses man



No glasses man



No glasses woman



−

+

=

Radford et al,  
ICLR 2016

Woman with glasses

