

Generative Adversarial Networks (GANs)

Adapted from material by Goodfellow, Binglin, Shashank, Bhargav

GANs

- **Generative**

- Learn a generative model

- **Adversarial**

- Trained in an adversarial setting

- **Networks**

- Use Deep Neural Networks

Why Generative Models?

- **Discriminative models:**

- Given an image \mathbf{X} , predict a label \mathbf{Y}
- Estimates $\mathbf{P}(\mathbf{Y}|\mathbf{X})$ Text

- **Discriminative models have several key limitations**

- Can't model $\mathbf{P}(\mathbf{X})$, i.e. the probability of seeing a certain image
- Thus, can't sample from $\mathbf{P}(\mathbf{X})$, i.e. **can't generate new images**

- **Generative models try to address these:**

- model $\mathbf{P}(\mathbf{X})$
- generate new data (e.g. images)

Magic of GANs...

Which one is Computer generated?



Magic of GANs...

User edits



Generated images



Adversarial Training

- Generator: generate fake samples, tries to fool the Discriminator
fake mean computer generated, not human generated
- Discriminator: tries to distinguish between real and fake samples
real as in human generated
- Train them against each other
- Repeat this and we get better Generator and Discriminator

GAN's Architecture

sampling process give to discrimination

discrimination dont know if its sample is real or not but it has to predict. binary classifier

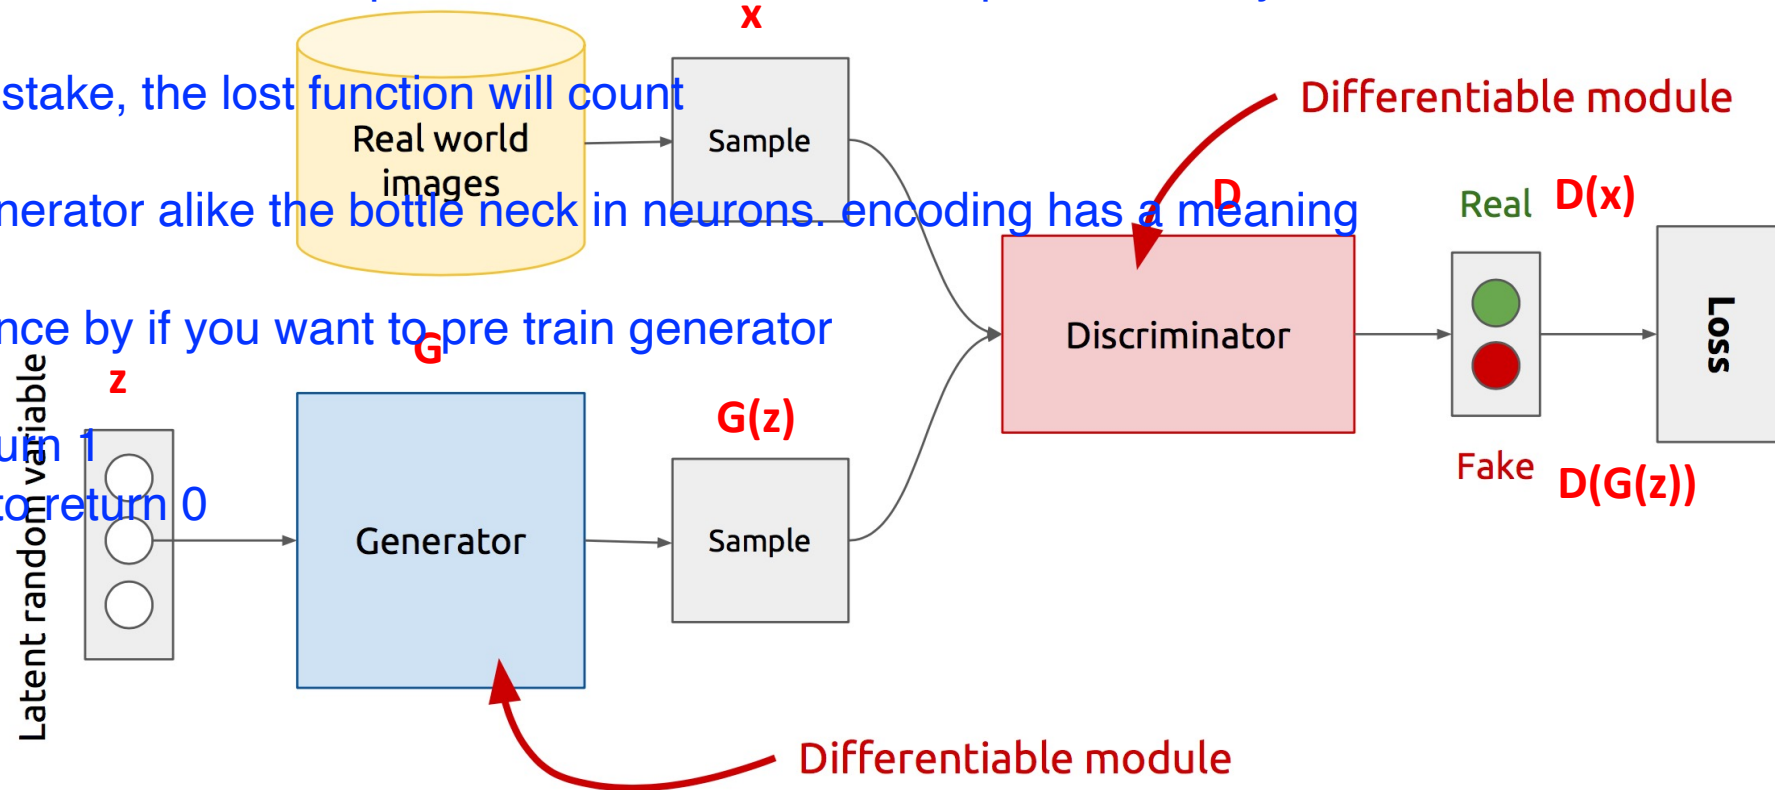
if it makes a mistake, the lost function will count

can think of generator alike the bottle neck in neurons. encoding has a meaning

archi not influence by if you want to pre train generator

if x , want to return 1

if z , then want to return 0



- Z is some random vector (Gaussian/Uniform).

- Z can be thought as the latent representation of the image.

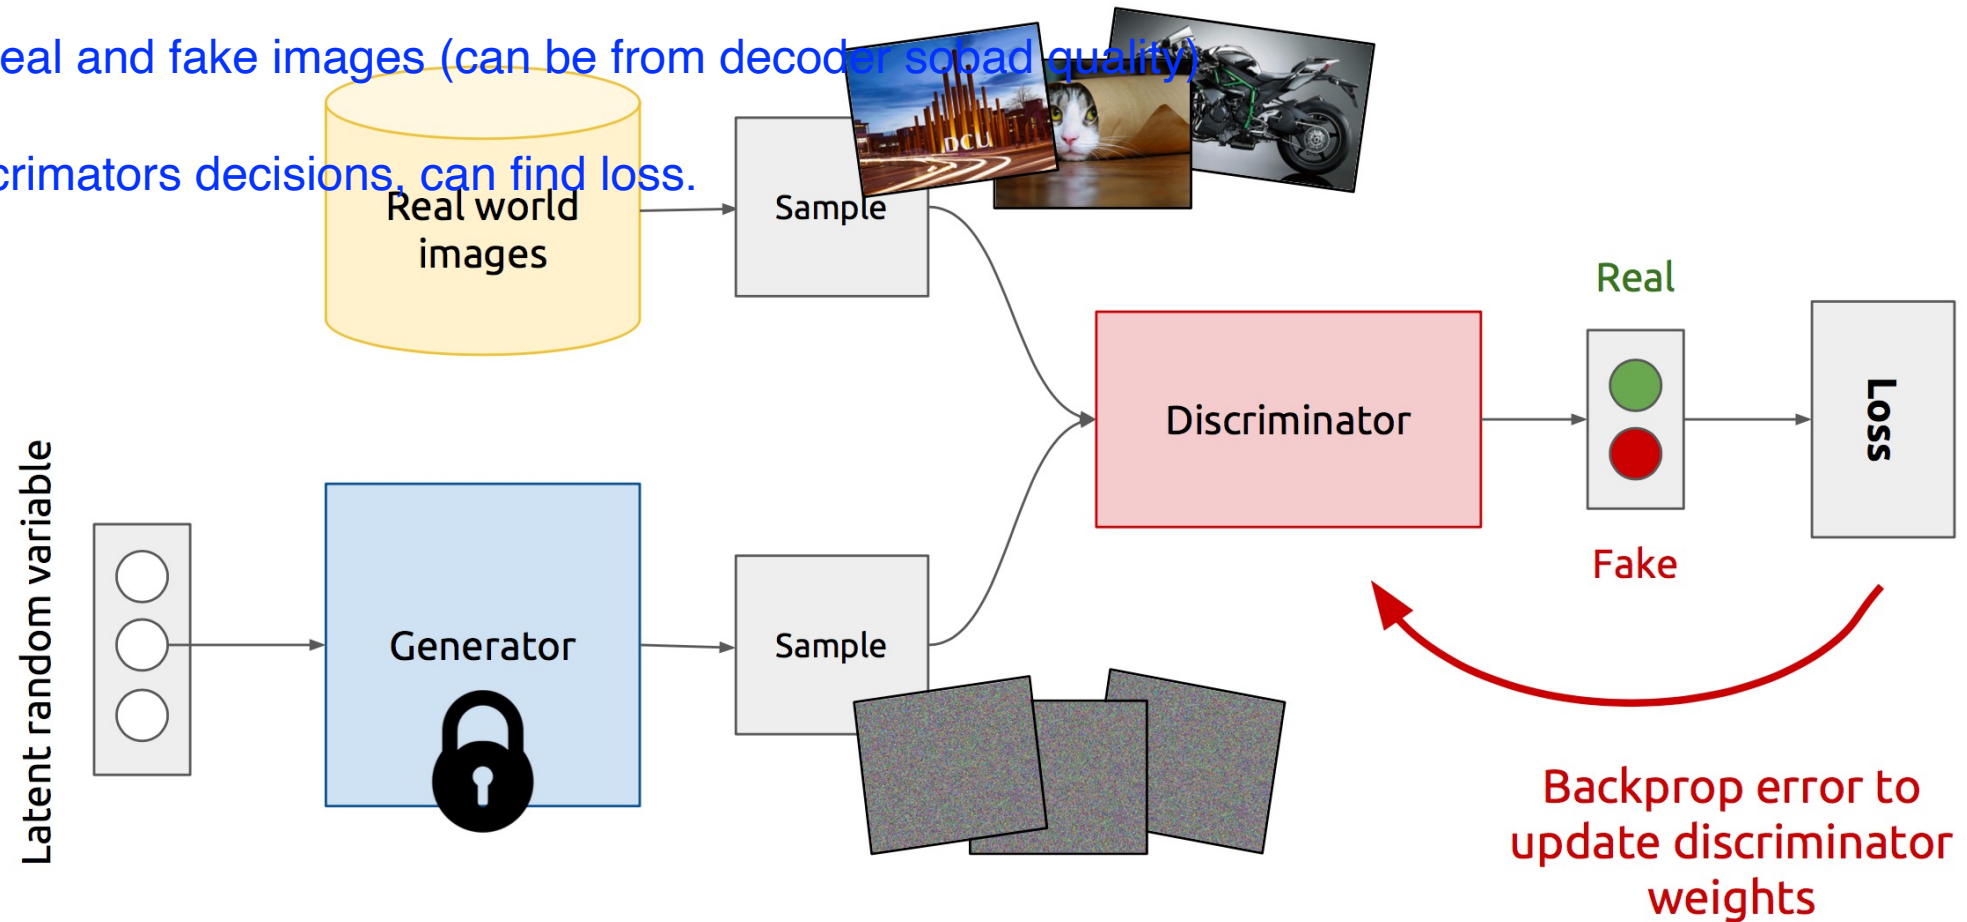
random in this latent space,
but has encoding,
has meaning

Training Discriminator

assume generator is lock

sample real and fake images (can be from decoder so bad quality)

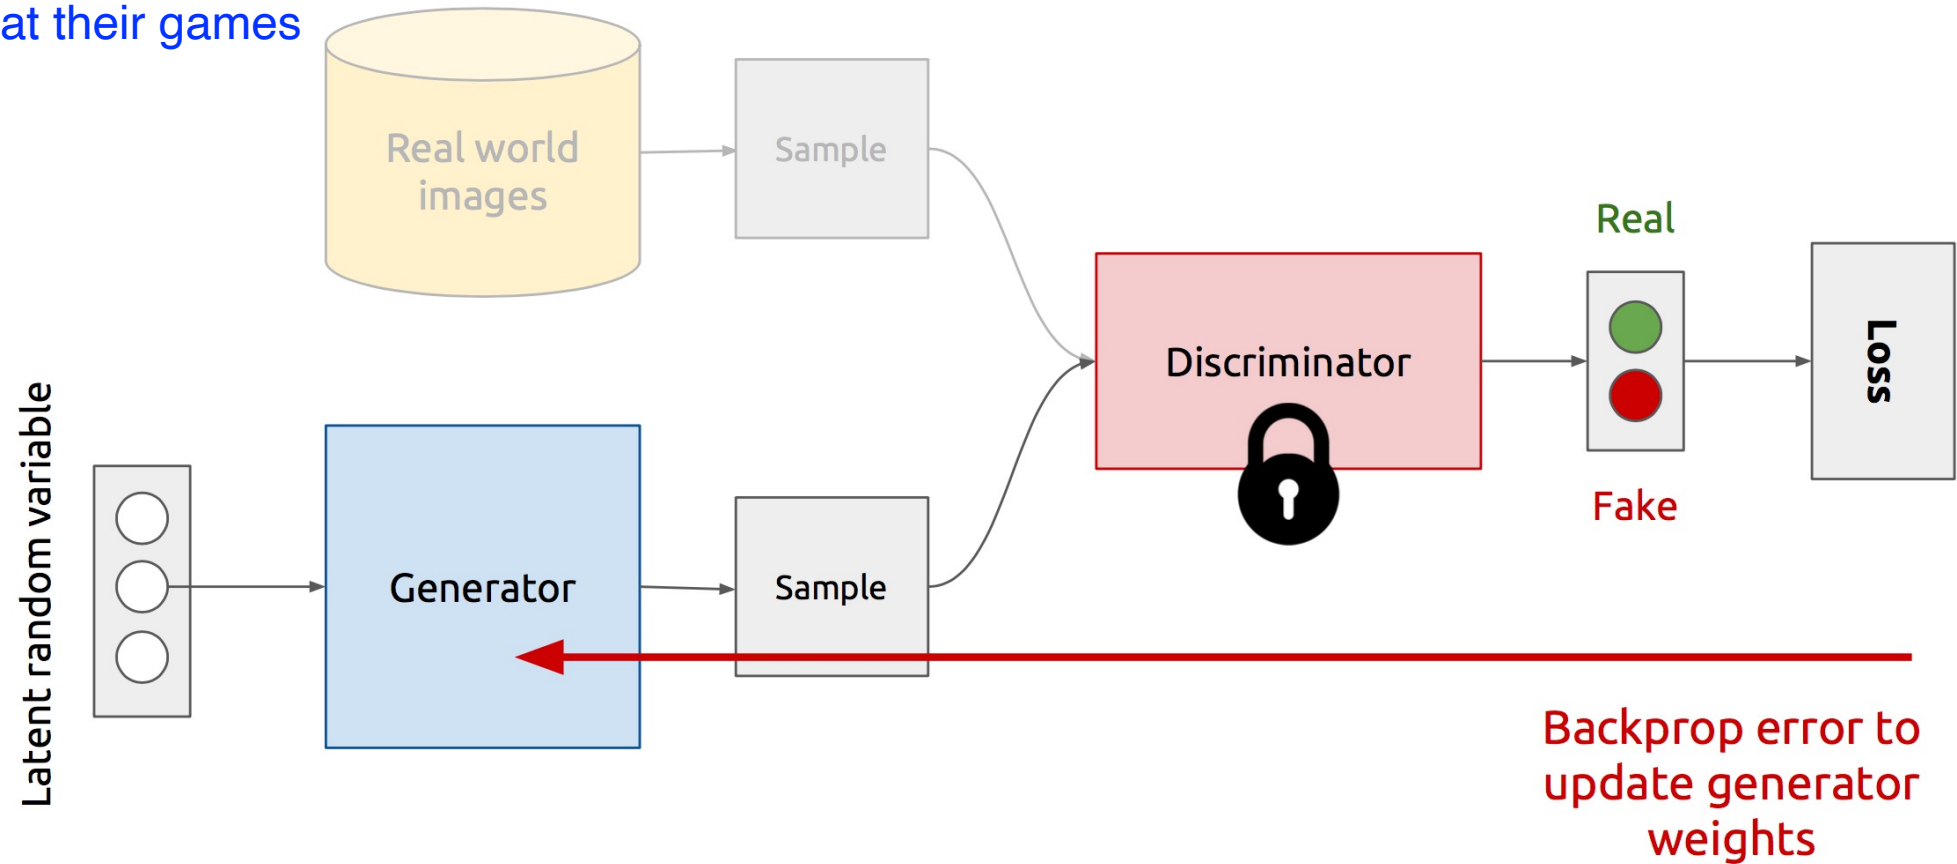
from discriminators decisions, can find loss.



Training Generator

training generator part can be updated by fixing discriminator

adjust weight so generator produce values that discriminator can make a mistake. repeat until both become good at their games



GAN's formulation

$$\min_G \max_D V(D, G)$$

- It is formulated as a **minimax game**, where:

- The Discriminator is trying to maximize its reward $V(D, G)$
- The Generator is trying to minimize Discriminator's reward (or maximize its loss)

G want to minimize reward for D
by producing realistic output

overall reward function

reward opposite of loss

$$V(D, G) = \mathbb{E}_{x \sim p(x)} [\log D(x)] + \mathbb{E}_{z \sim q(z)} [\log(1 - D(G(z)))]$$

if x coming from x, then $E[\log D(x)]$ want to return 1

on other hand, if z is latent variable from the fake, D decide real or fake, want it to be able to determine

- The Nash equilibrium of this particular game is achieved at:

- $P_{data}(x) = P_{gen}(x) \quad \forall x$
- $D(x) = \frac{1}{2} \quad \forall x$

as they become better at their games, will reach optimal players
so D output 0.5, cant distinguish 50% 50% chance

Algorithm 1 Minibatch stochastic gradient descent training of generative adversarial nets. The number of steps to apply to the discriminator, k , is a hyperparameter. We used $k = 1$, the least expensive option, in our experiments.

for number of training iterations **do**

for k steps **do**

- Sample minibatch of m noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from noise prior $p_g(z)$.
- Sample minibatch of m examples $\{x^{(1)}, \dots, x^{(m)}\}$ from data generating distribution $p_{\text{data}}(x)$.
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[\log D(x^{(i)}) + \log (1 - D(G(z^{(i)}))) \right].$$

end for

- Sample minibatch of m noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from noise prior $p_g(z)$.
- Update the generator by descending its stochastic gradient:

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log (1 - D(G(z^{(i)}))).$$

end for

The gradient-based updates can use any standard gradient-based learning rule. We used momentum in our experiments.


**Discriminator
updates**

**Generator
updates**

Faces



CIFAR

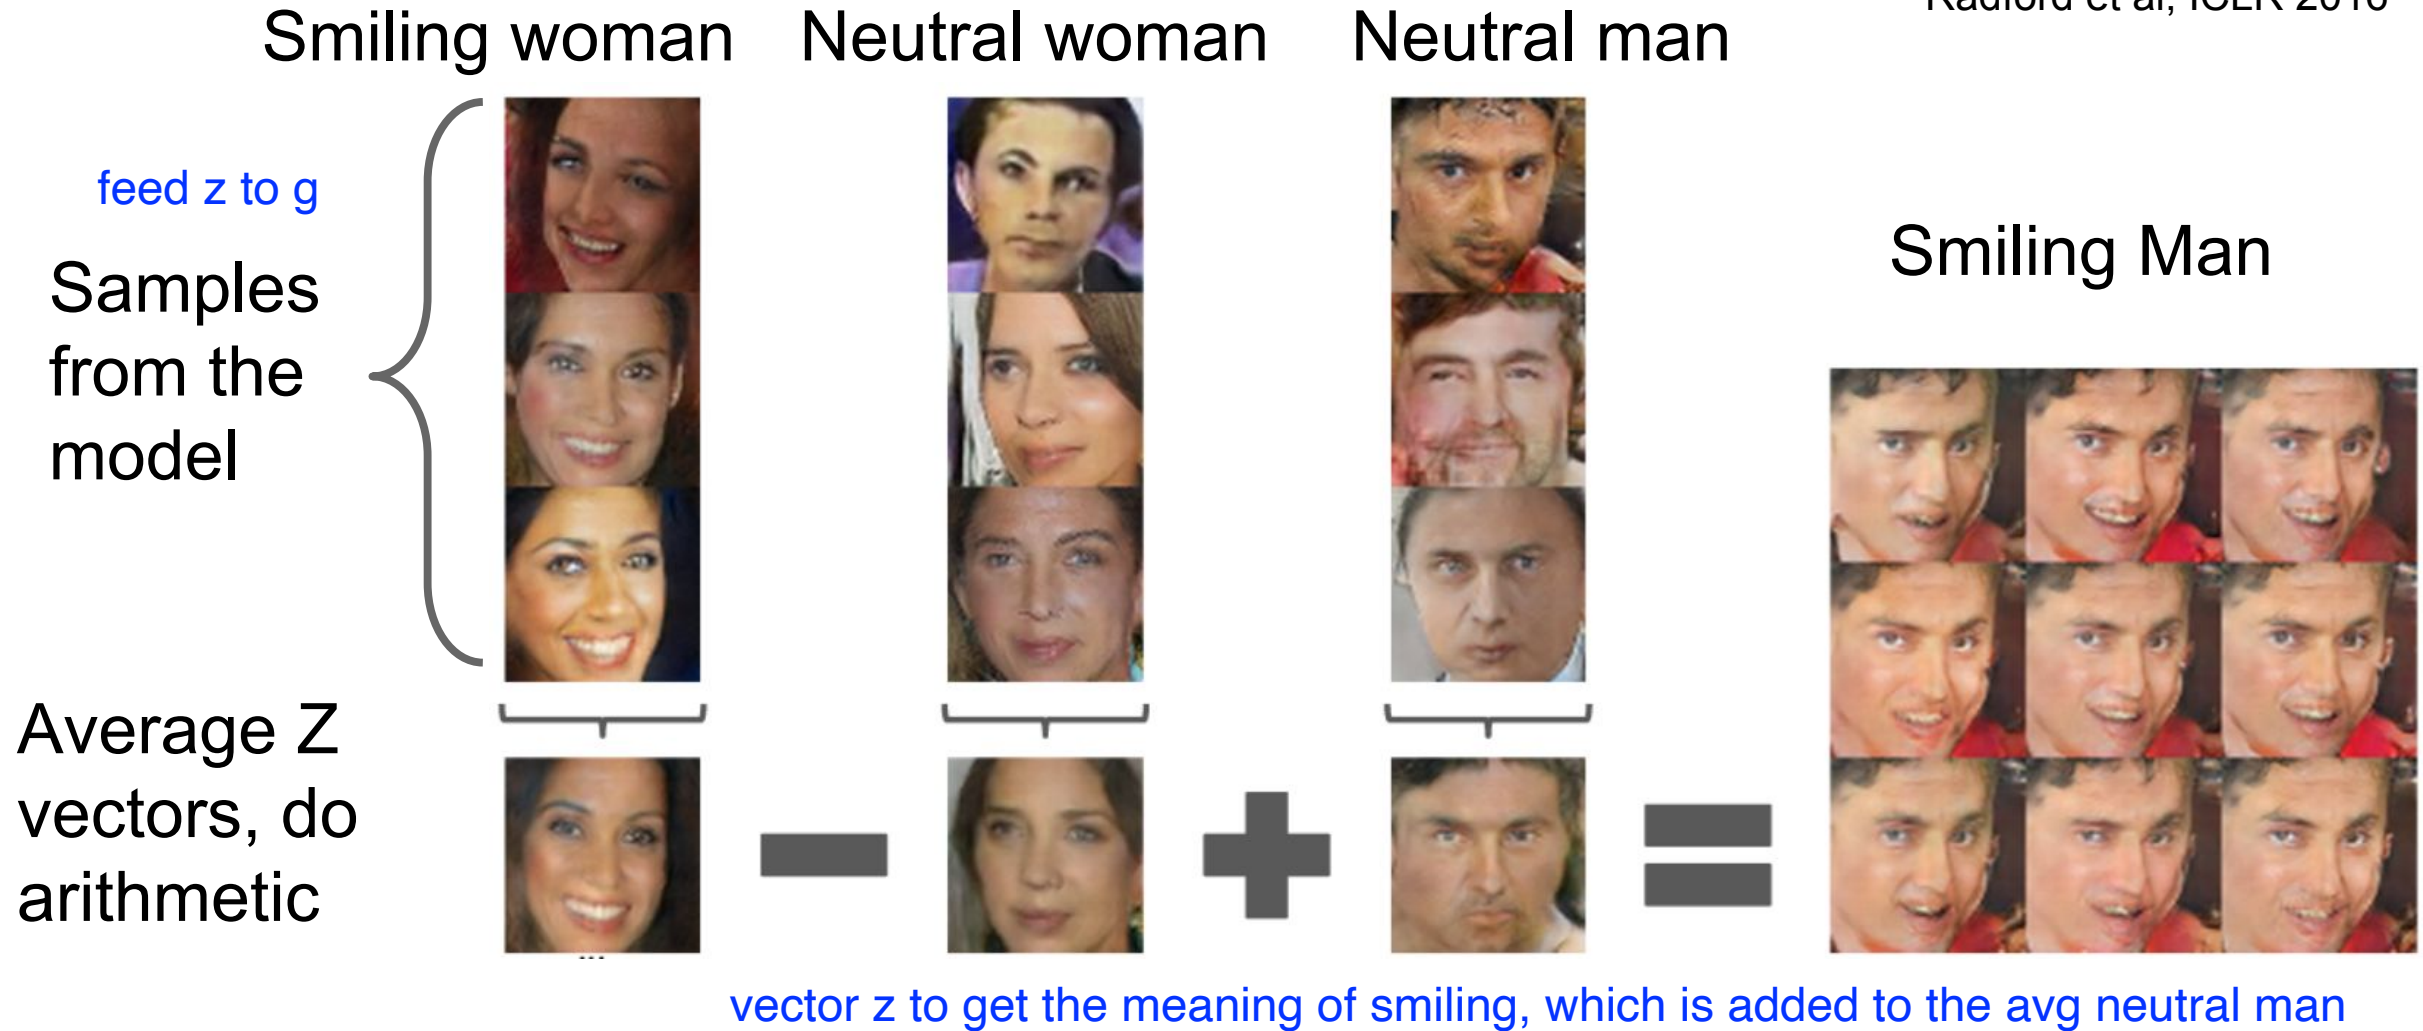


DCGAN: Bedroom images



Generative Adversarial Nets: Interpretable Vector Math

Radford et al, ICLR 2016



Generative Adversarial Nets: Interpretable Vector Math

Glasses man



No glasses man



No glasses woman



Radford et al,
ICLR 2016

Woman with glasses



−

+

=