

Supervised learning. Want to look at number of examples.

Training data being cartesian product of $X \times Y$

Reducing to a search where goal is to find h

1 Function approximation

Consider a sample $D \subset \mathcal{X} \times \mathcal{Y}$ where the data is generated by a process, and \mathcal{X} and \mathcal{Y} can be seen as input and output (or target) respectively:

$$D = \{(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)\}$$

The goal of function approximation is to find (or learn) a function $h \in \mathcal{H}$ where $h: \mathcal{X} \rightarrow \mathcal{Y}$ such that given an x , it “well-approximates” the corresponding y value.

Depending on whether the target is **continuous** or **discrete**, the task is called **regression** or **classification** respectively.

Examples include recognition of handwritten digits, prediction of house prices in Christchurch, and prediction of Students’ performance in a course.

There are different ways of approaching this problem depending on what assumptions are made. For example y_i can be assumed to be generated by

- $y_i = f(x_i)$ • **a function** (i.e. a deterministic/exact process); or **process generating a deterministic output that doesn't vary. Same input -> same output**
- $y_i \sim F_{\theta|x}(x_i)$ • **a probability distribution** (i.e. a stochastic process).

2 A logical approach to binary classification

In binary classification \mathcal{Y} is $\{True, False\}$ (or $\{0, 1\}$, or $\{positive, negative\}$). All classification problems can be cast as one or more binary classification problems.

eg house prices, just looking at area as input,

The target is $y = f(x)$ where f is a deterministic (but unknown) Boolean-valued function: $f: \mathcal{X} \rightarrow \{True, False\}$. In other words, f is an unknown *predicate*.

same input can get diff outputs.

When a sample is given, it is assumed that the sample is generated as:

so need to take in more input and probabilistic approach

$$D = \{(x_1, f(x_1)), (x_2, f(x_2)), \dots, (x_n, f(x_n))\}$$

Let \mathcal{F} be the set of all predicates (or functions) that can be defined over \mathcal{X} . In this setting we have:

Don't go to all set of \mathcal{F} , as \mathcal{F} is large
Can't just get f as f is unknown

$$\mathcal{H}, \text{ hypothesis } \mathcal{H} \subseteq \mathcal{F}$$

$$2^{|\mathcal{X}|}$$

Question What is the cardinality (size) of \mathcal{F} if \mathcal{X} is finite?

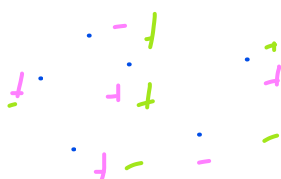
Answer: $2^{(\text{size of } \mathcal{X})}$

\mathcal{F} is set of predicates, so binary true or false

\mathcal{X} in our universe is all possible inputs

eg one function in green, another in pink

efficiency and size will become a consuming problem



think of functions as its input and output.

2.1 Partial order of predicates Predicate test if something holds or not

A relation R on a set A is a partial order if it is reflexive, transitive, and anti-symmetric.

Reflexive,
 $aRa \quad aA.$

Antisymmetric
 $aRb \text{ and } bRa$
 $a = b.$

Transitive
 $aRb \text{ and } bRc$
 $aRc.$

We define the binary relation *less general or equal* over \mathcal{F}^2 (or \mathcal{H}^2) as the following:



if first one is true, then second one also true,
but not necessarily the other way around.

$$h_a \leq_g h_b \quad \equiv \quad \forall x \in \mathcal{X} \quad h_a(x) \implies h_b(x)$$

The relation *less general or equal* has the following properties:

- reflexive: $\forall h \in \mathcal{H} : \quad h \leq_g h$
- anti-symmetric: $\forall h_a, h_b \in \mathcal{H} : \quad ((h_a \leq_g h_b \wedge h_b \leq_g h_a) \implies h_a = h_b)$
- transitive: $\forall h_a, h_b, h_c \in \mathcal{H} : \quad ((h_a \leq_g h_b \wedge h_b \leq_g h_c) \implies h_a \leq_g h_c)$

Therefore the relation establishes a partial order over \mathcal{F} (or \mathcal{H}).



contrast with total order

We define the *support* of a hypothesis to be the set of elements in \mathcal{X} that are accepted by the hypothesis. In other words, $supp : \mathcal{H} \rightarrow 2^{\mathcal{X}}$ and

$$supp(h) = \{x \in \mathcal{X} : h(x)\}$$

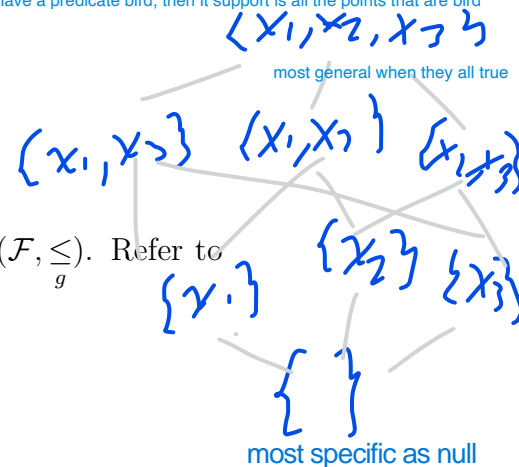
so if have a predicate bird, then it support is all the points that are bird

As a result we have

Support of h_a is a subset of h_b

$$h_a \leq_g h_b \quad \implies \quad supp(h_a) \subseteq supp(h_b)$$

Example suppose $\mathcal{X} = \{x_1, x_2, x_3\}$. Draw the lattice of (\mathcal{F}, \leq_g) . Refer to the elements of \mathcal{F} by their support sets.



2.2 Consistency and learning

2.2.1 Consistency

For an $h \in \mathcal{H}$, we define consistency as the following:

$$consistent(h, D) \quad \equiv \quad \forall (x, y) \in D : h(x) = y$$

Example Write a Python function `consistent(h, D)`.

2

Answer:

```
def consistent ( h, D )
    return all( h(x) == y for x, y in D )
```

$$2^{(|X| - |D|)}$$

D in reality could be inconsistent, but in this example we assume not

if D was size of X, if D was empty. are the extremes

Question How many functions in \mathcal{F} can be consistent with D? Half?

Answer: $2^{|X| - \text{cardinality of } D}$

2.2.2 Learning one hypothesis

One (inefficient) way of learning is by exhaustively searching \mathcal{H} .

finding a h from H

Example Write a Python function `learn_a_hypothesis_by_enumeration(H, D)`. parameters are H and D, for consistent it was h, D

Is it possible for the returned classifier to misclassify a new input? as only concern with consistency with D so can misclassify if data not seen before, not from D. a real concern

Often the elements of \mathcal{H} have the same type and encoding. If \mathcal{R} is the set of all codes for \mathcal{H} , there is a function to map a code to a hypothesis: Use R as set of codes expecting H to be passed as already made set

$\text{decode} : \mathcal{R} \rightarrow \mathcal{H}$

decode is a function that takes a code (an argument) in R, and return function in H.

Often there is no bijection between \mathcal{H} and \mathcal{R} .

-> means it is a mapping from a r in R to a h in H?

The elements of \mathcal{H} are rarely generated explicitly; they come in to existence as necessary. The search is performed in \mathcal{R} . This is similar to the way in which vertices and edges in implicit graphs come to existence.

d for the weight, and 1 for the bias

$$R = \mathbb{R}^{d+1}$$

Example Give \mathcal{R} for the following hypothesis classes:

- \mathcal{H} is the set of perceptrons over \mathbb{R}^d represented by weights and bias.
- \mathcal{H} is the set of naive Bayes networks where the objects are described by d Boolean input variables and the class label is binary.

Decode take the weight and bias and return true or false

Decode is part of the design

3 The version space

The version space is defined as where function is deterministic


$$VS_{\mathcal{H}, D} = \{h \in \mathcal{H} : \text{consistent}(h, D)\}$$

Example Write a Python function `learn_all_hypotheses(H, D)` that finds the version space by enumeration.

all elements that are consistent with H will be return

H is always a subset of F


3



diff weights can produce same perceptron. sometimes theres lots of effort to remove this redundancy

$$[0, 1]^{d+1}$$

integral, not a list



consider definition of version space depend on D, H, what set of hypothesis and data you have. need D because talking about consistency, need H as talking about functions


partial order with less general and equal to hierarchy graph where there is none less than or equal to general in the same hierarchy row.

can have two bounds

as general as you can be and still be consistent with your data. Highlighted is the version space that is consistent with your data that is the most specific is S, and G is most general that can still be specific

Example The entire version space can be used for prediction.

most general



most specific

```
def predict(VS, x):
    if len(VS) == 0:
        raise ValueError("No hypothesis!")
    if all(h(x) for h in VS):
        return "All positive"
    elif not any(h(x) for h in VS):
        return "All negative"
    else:
        return "positive count:{}, negative count:{}".format(# complete)
```

may or may not be correct as x could be out of D. if it in D, and do find a hypothesis, then they are going to agree with each other,.

if not in D, they will disagree,

3.1 Specifying VS with its boundaries

We define two sets:

$$S_{H,D} = \{h \in \mathcal{H} : \text{consistent}(h, D) \wedge \forall h' \in \mathcal{H} : (\text{consistent}(h', D) \implies \neg(h'_g < h))\}$$

$$G_{H,D} = \{h \in \mathcal{H} : \text{consistent}(h, D) \wedge \forall h' \in \mathcal{H} : (\text{consistent}(h', D) \implies \neg(h < h'_g))\}$$

See Figure 1 for an illustration of the version space and the sets S and G .

Theorem 1.

$$\forall h \in \mathcal{H} : (h \in VS_{H,D} \Leftrightarrow \exists h_s \in S_{H,D}, \exists h_g \in G_{H,D} : h_s \leq_g h \leq_g h_g)$$

D is number of positive and negative training examples.

imagine each point has a coordinate of x and y

hypothesis space is all rectangles.

if something falls inside rectangle, it will return true.

cannot have a rectangle smaller than SB that can be consistent with the data.

GB is the biggest it can be before it will return positive for something that's actually negative

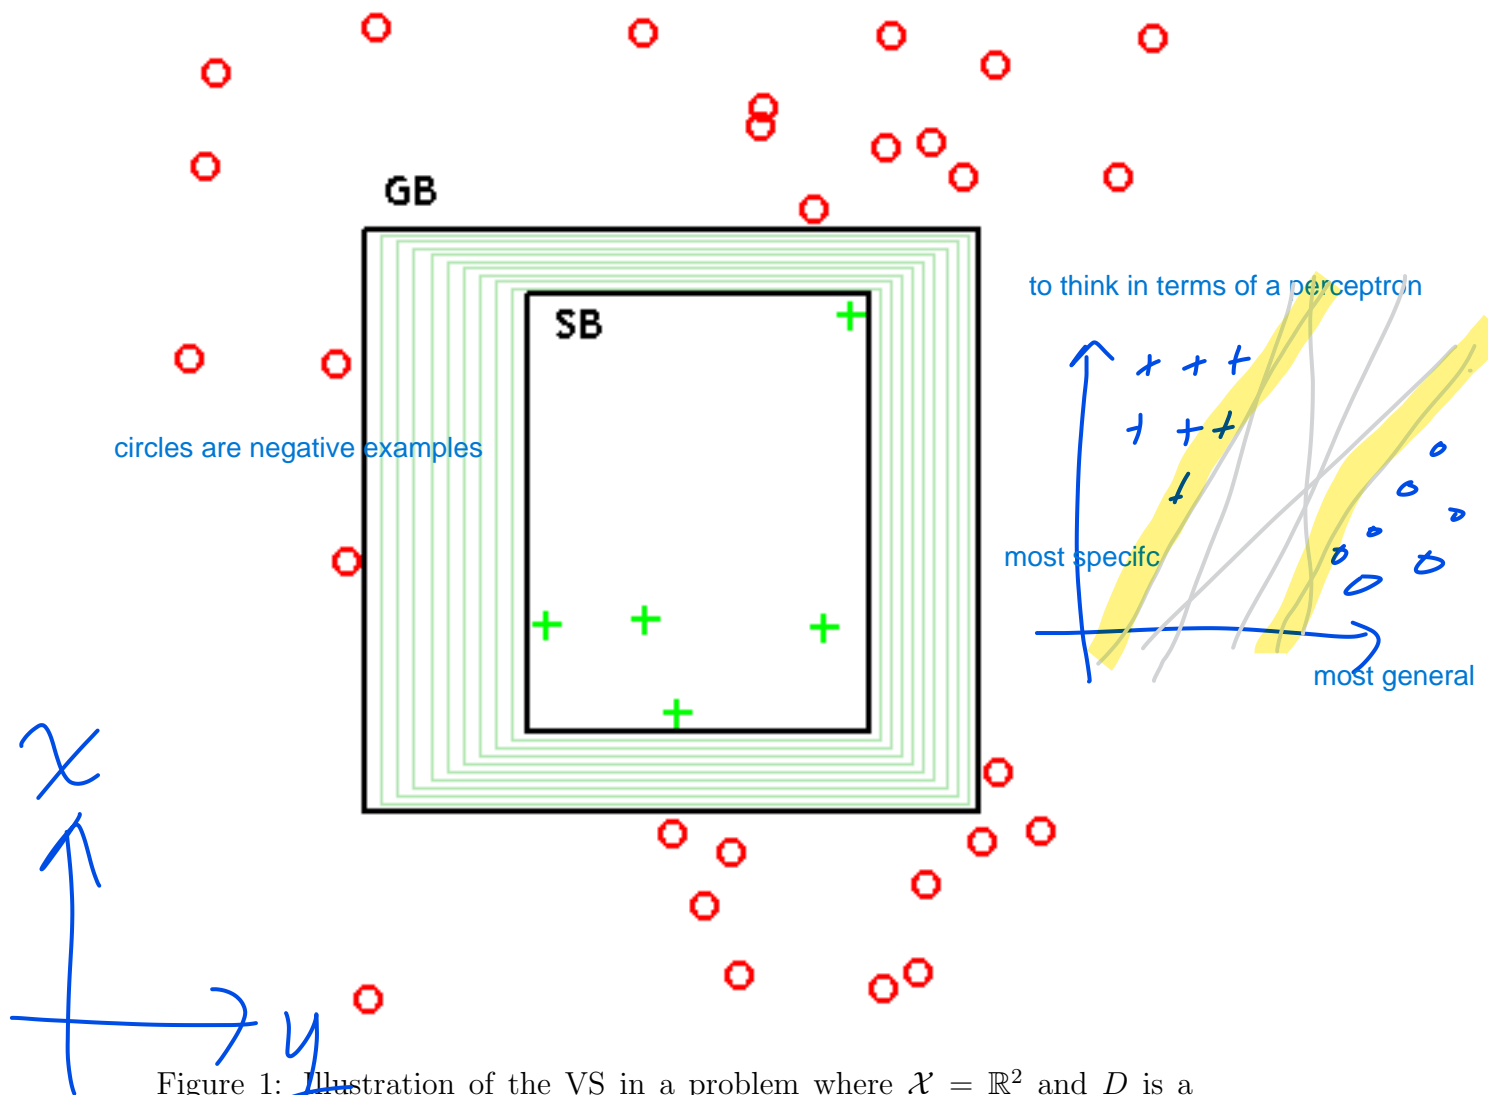


Figure 1: Illustration of the VS in a problem where $\mathcal{X} = \mathbb{R}^2$ and D is a collection of positive points (depicted by '+') and negative points (depicted by 'o'). \mathcal{H} is the set of all rectangles in \mathbb{R}^2 where a point inside a rectangle will be predicted as positive. The version space can be specified by two sets S and G , each having one member in this example.

the idea is to learn a VS not by enumerating over all the elements in H as this is expensive, but instead to just look at S and G as they implicitly define the version space.

slowly moving S and G tighter to be more specific while still being most general as you iterate to see all the data

3.2 Candidate Elimination Algorithm (CEA)

The idea is to incrementally update S and G instead of generating VS by filtering through \mathcal{H}

Initialize G to the set of most-general hypotheses in H

Initialize S to the set of most-specific hypotheses in H

For each training example, d , do:

 If d is a positive example then:

 Remove from G any hypotheses that do not match d

 For each hypothesis s in S that does not match d

 Remove s from S

 Add to S all minimal generalizations, h , of s such that:

 1) h matches d

 2) some member of G is more general than h *** more general or equal

 Remove from S any h that is more general than another hypothesis in S

 If d is a negative example then:

 Remove from S any hypotheses that match d

 For each hypothesis g in G that matches d

 Remove g from G

 Add to G all minimal specializations, h , of g such that:

 1) h does not match d

 2) some member of S is more specific than h ** more specific or equal

 Remove from G any h that is more specific than another hypothesis in G

4 Example Problem and Representation

4.1 Representation: Conjunction of Constraints

If $\mathcal{X} = \mathcal{X}_1 \times \mathcal{X}_2 \times \dots \times \mathcal{X}_d$, then we define the representation of *conjunction of constraints* (i.e. the set of all codes) to be

1-d features in total. position 1 is everything in domain of 1 union with wildcard and null

$$\mathcal{R} = \prod_{i=1}^d (\mathcal{X}_i \cup \{?, \emptyset\})$$

Each code is a tuple of the form (c_1, c_2, \dots, c_d) where c_i is a constraint which is either:

- ? (a wild card): meaning that the value of the i -th feature can be anything;
- or

- a specific value from the domain of the i -th feature: meaning that the value of the i -th feature must be exactly the given specific value; or
- \emptyset : meaning that no value for the i -th feature is acceptable.

Given a tuple of constraints, the corresponding hypothesis is the conjunction of all the d constraints. More formally let

$$satisfied(c_i, x_i) \mapsto c_i \neq \emptyset \wedge (c_i = ? \vee c_i = x_i)$$

then

must be a wild card or exactly that value

$$decode(c)(x) \mapsto satisfied(c_1, x_1) \wedge satisfied(c_2, x_2) \wedge \dots \wedge satisfied(c_d, x_d)$$

take a constraint that behave in this way. last question in quiz

4.2 Example problem

This problem is about when a certain sport is enjoyable. The input space is defined by 6 attributes: $\mathcal{X} = Sky \times Temp \times Humidity \times Wind \times Water \times Forecast$. The target is a Boolean: $\mathcal{Y} = \{Yes, No\}$. The domain of the variables are as the following:

- $Sky = \{\text{Sunny, Cloudy, Rainy}\}$
- $Temp = \{\text{Warm, Cold}\}$
- $Humidity = \{\text{Normal, High}\}$
- $Wind = \{\text{Strong, Weak}\}$
- $Water = \{\text{Warm, Cool}\}$
- $Forecast = \{\text{Same, Change}\}$

Questions

1. What is the cardinality of \mathcal{X} ? $3 * 2 * 2 * 2 * 2 * 2 = 3 * 2^5 = 96$
2. What is the cardinality of \mathcal{F} ? 2^96
3. If we use the conjunction of constraints as our hypothesis class, what is the size of R ? What is the size of \mathcal{H} ?

$$|R| = (3+2) * (2+2)^5$$

if you have a null, then entire hypothesis is null.
 $|H| = 1 + (3+1) * (2+1)^5$

$|H|$ is much smaller compare to 2^96
 $|H|$ is almost always smaller than $|R|$

best case $|H| = |R|$

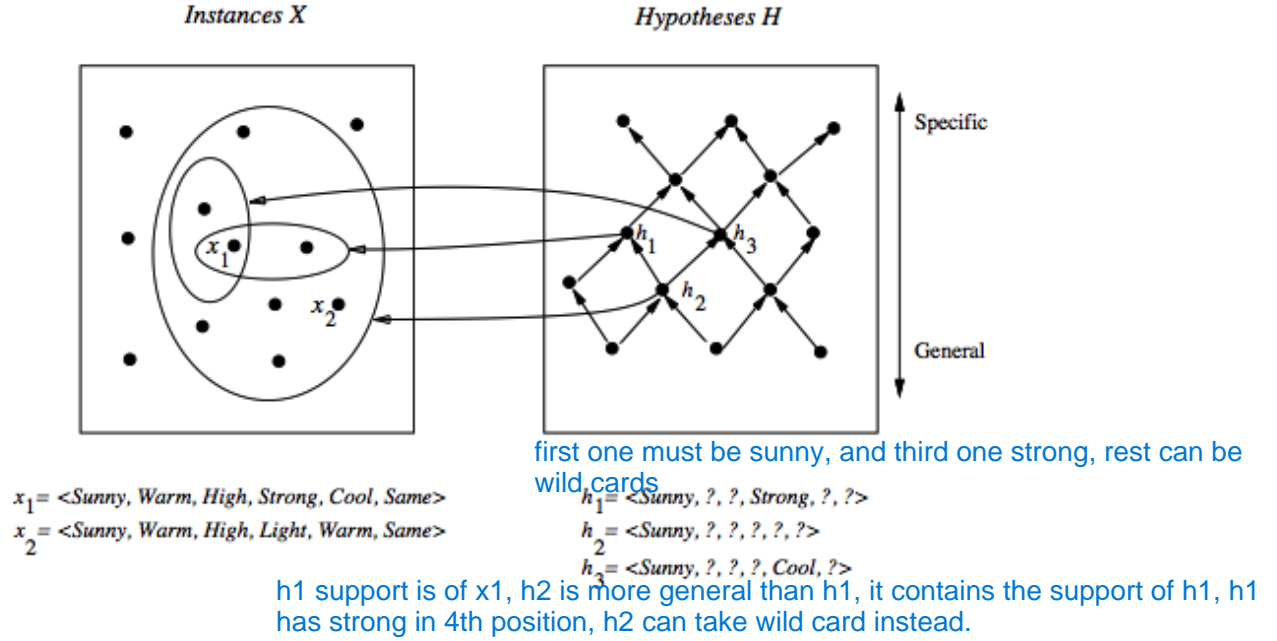


Figure 2: The relationship between a few hypotheses each expressed as conjunction of constraints and a few points in the input space.

4.3 Applying CEA

Consider the following sample D :

Sky	Temp	Humid	Wind	Water	Forecast	Enjoy Sport
Sunny	Warm	Normal	Strong	Warm	Same	Yes
Sunny	Warm	High	Strong	Warm	Same	Yes
Rainy	Cold	High	Strong	Warm	Change	No
Sunny	Warm	High	Strong	Cool	Change	Yes

Figure 3 shows the trace of CEA using the conjunction of constraints and the above D .

The resulting VS is specified with only S and G . See Figure 4.

If S and G become empty sets after the execution of the algorithm, then there is no hypothesis (predicate) in \mathcal{H} that is consistent with the data (i.e. no learning).

from the bottom, the beginning G node has wildcard for everything. hypothesi that accept everything and return true for everything.

top S0, ust need one null to return null

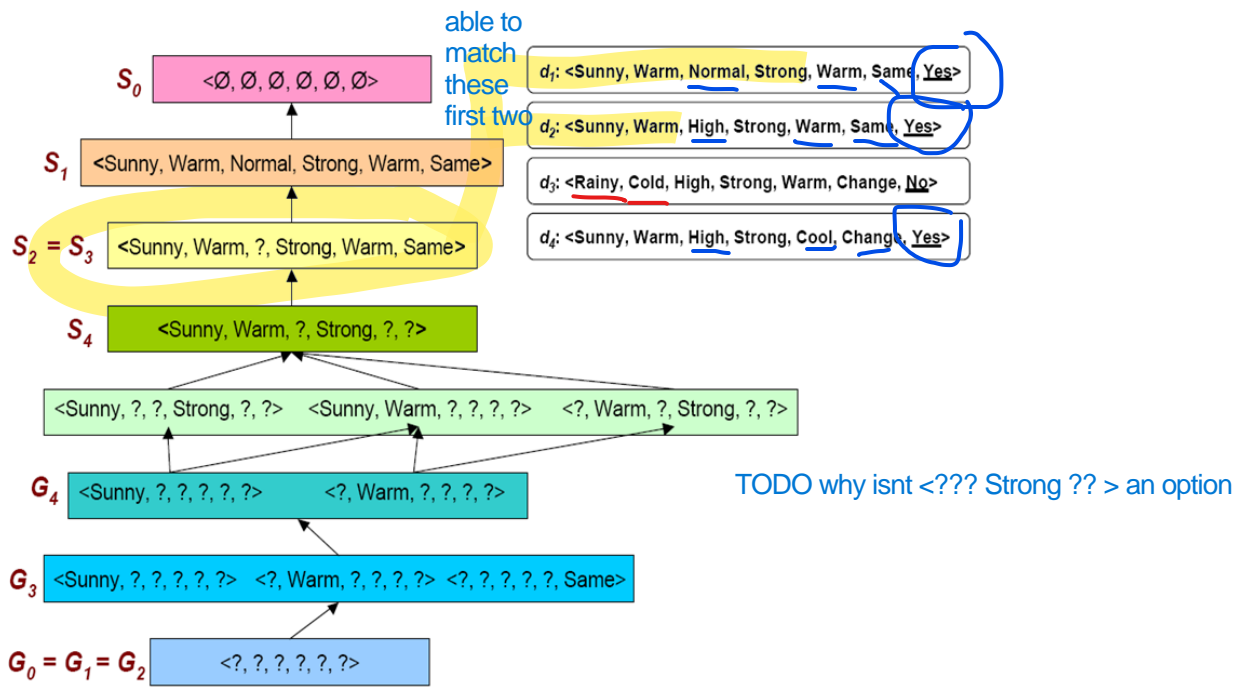


Figure 3: The trace of CEA using conjunction of constraints and four training data points.

Making predictions

Using the produced VS, determine how the following cases are classified:

- (Sunny, Warm, Normal, Strong, Cool, Change)
- (Rainy, Cool, Normal, Light, Warm, Same)
- (Sunny, Warm, Normal, Light, Warm, Same) Y?
- (Sunny, Cold, Normal, Strong, Warm, Same) Y?

4.4 Language (representation) limitations

What if D is the following set?

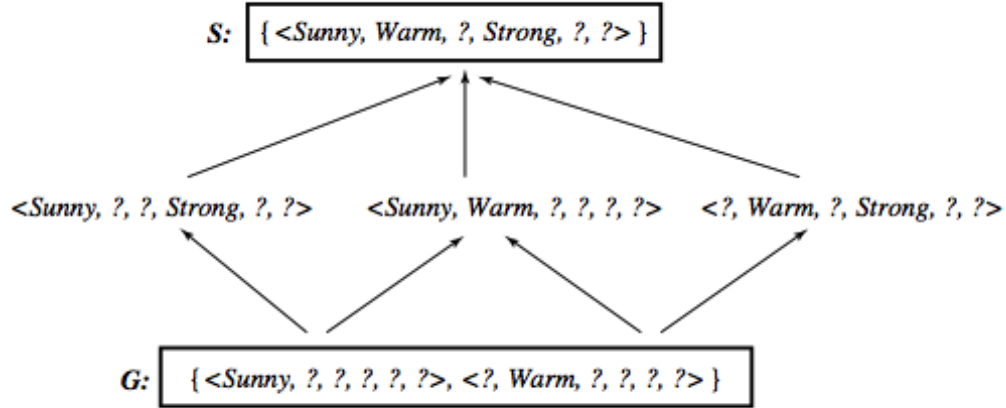


Figure 4: The resulting VS which can be specified using only S and G . If necessary, S and G can be used to generate the rest of the hypotheses in the VS.

Sky	Temp	Humid	Wind	Water	Forecast	Enjoy Sport
Sunny	Warm	Normal	Strong	Cool	Change	Yes
Cloudy	Warm	Normal	Strong	Cool	Change	Yes
Rainy	Warm	Normal	Strong	Cool	Change	No

- No hypothesis in the form of conjunction of constraints is consistent with data. because if make first position a wild card it will accept everything, if makee cloudy then it will be yes for second row but no for first row
- The hypothesis $(?, \text{Warm}, \text{Normal}, \text{Strong}, \text{Cool}, \text{Change})$ is too general.
- Target concept is in a different hypothesis space. For example in space where hypotheses can have disjunction: $(\text{Sky} == \text{"Sunny"} \text{ or } \text{Sky} == \text{"Cloudy"})$

5 Inductive Bias

Any means by which a learning system is limited to choose from only a limited set of functions (hypotheses) is *inductive bias*. There are two major forms of bias:

lots of freedom to learning $H = F$ then you just ;not learning more than what we observe where in this course you would like to generalize, where this is not possible without bias

- Language bias: the representation is such that \mathcal{H} contains only certain functions.
- Search bias: The search algorithm embodies a preference for certain functions and does not consider all options.

both 2^X

allowing all logical formulas, able to express any function so any F is in H , H is size of F

5.1 Extending the hypothesis space

We can allow all logical formulas and therefore be able to express any function. What will be the size of \mathcal{H} in this case? What can we learn in this case?

For example, consider $D = \{(x_1, True), (x_2, True), (x_3, True), (x_4, False), (x_5, False)\}$. After running CEA we have:

- $S = \{(x_1 \vee x_2 \vee x_3)\}$ true if exactly x_1 , exactly x_2 , exactly x_3
- $G = \{\neg(x_4 \vee x_5)\}$ if its not x_4 or not x_5 , it'll return true

- We can only classify precisely examples already seen!
- Take a majority vote?
- For any hypothesis h in VS that classifies an unseen x as positive, there is a complementary hypothesis that classifies x as negative.
- Unseen instances are classified positive (and negative) by exactly half of the VS .

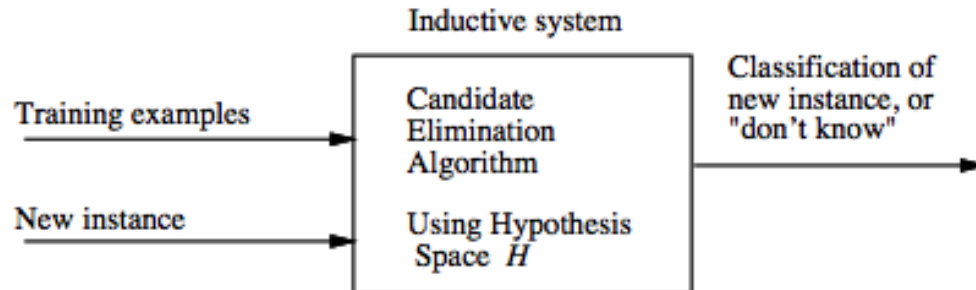
5.2 No inductive inference without a bias

A learner that makes no a priori assumptions regarding the identity of the target concept, has no rational basis for classifying unseen instances.

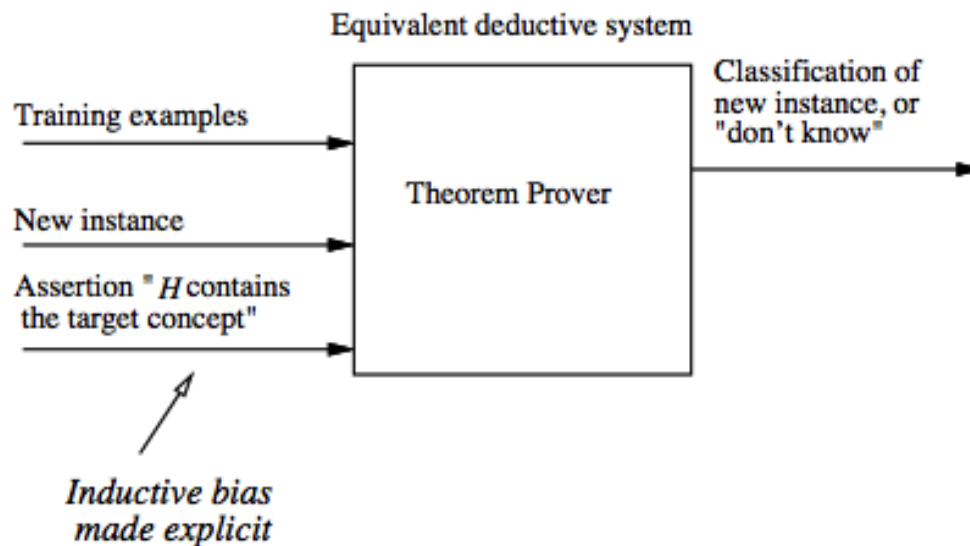
6 Inductive learning as deductive process

Induction is inferring general rules and theories from specific empirical data (this is a broader concept than logical/mathematical induction). The inductive learning process described earlier is depicted in the following diagram.

if you manage to describe your bias as a set of assumptions then can in fact use a deductive process (eg. prolog) to actually do learning to describe functions too - not something we need to do in practice.



Deduction is inferring sound specific conclusions from general rules (axioms) and specific facts. In a logical approach to learning, the inductive process can be achieved using deduction.



7 A probabilistic approach

7.1 Verifying a hypothesis

Given a hypothesis, we want to know the probability of the hypothesis making an error. Let's define a Bernoulli random variable that is 1 when an error happens:

indicator function return 0 or 1

name of function is 1

subscription describe this is about $h \neq f$

return 1 is there is an error for a given input (x)

$$\mathbb{1}_{h \neq f}(x) = \begin{cases} 1, & h(x) \neq f(x) \\ 0, & \text{otherwise.} \end{cases}$$

We define the following quantity as the error of h :

error of h is prob that h does not agree with f = expected value of the indicator
 $E(h) = \mathbb{P}_{\mathcal{X}}[h(x) \neq f(x)] = \mathbb{E}_{\mathcal{X}}[\mathbb{1}_{h \neq f}(x)]$,

where $\mathbb{P}_{\mathcal{X}}[h(x) \neq f(x)]$ denotes the probability of the event $h(x) \neq f(x)$, and $\mathbb{E}_{\mathcal{X}}[\mathbb{1}_{h \neq f}(x)]$ denotes the expected value of the random variable $\mathbb{1}_{h \neq f}(x)$.

For the following reasons, we cannot directly calculate the above quantity:

- Although in this framework f is fixed, it is unknown to us;
- We do not have access to the entire input space \mathcal{X} , but only a sample D from it.

We can however measure the (empirical) error over the sample:

$$E_D(h) = \mathbb{P}_D[h(x) \neq f(x)] = \mathbb{E}_D[\mathbb{1}_{h \neq f}(x)] = \frac{1}{|D|} \sum_{x \in D} \mathbb{1}_{h \neq f}(x) .$$

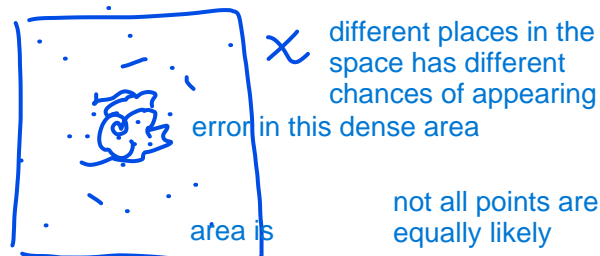
7.2 Hoeffding's inequality

For $\epsilon > 0$, the probability that the empirical error is off (the actual error) by more than ϵ is bounded from above by a decreasing function of $|D|$:

$$\mathbb{P}[|E(h) - E_D(h)| > \epsilon] \leq 2e^{-2\epsilon^2|D|}$$

Remarks:

- The sample points must be generated independently.
- The underlying distribution (process) must be fixed.
- The key factor in the bound is $|D|$.



f is true
 h is predictions

not all points are equally likely

more likely distributed in dense, and less likely in corners

imagine $f(x)$ in dense areas, if there's a h that makes error in corners but is good in center, then that good h

on other hand, h with lots of error in dense area but correct in corner then not a good h

in way weighing error in high Likelihood error, expected value,

upper bound in this form is a negative value, and there's two cases, where it is size of D . as D is larger (so more samples) then bound will approach zero and be smaller

make it general, but general H can work against you, in putting bound in your direction

how to determine size of H when working in the numeric space

use trick to look at it computationally, bit, floating point number as unique number, size of r, and get a value for H
 pure mathematical sense, VC notion. would make H very large.

16 bit representation of numbers

7.3 Learning a hypothesis

When learning, you are no longer *given* a hypothesis but rather *choose* one from \mathcal{H} . For example, you may pick a hypothesis h^* that has the lowest empirical error on the training data:

could have H that only agree with some data so it not really consistent

$$h^* = \operatorname{argmin}_{h \in \mathcal{H}} E_D(h)$$

arg is returning something from the h. using value of the error to pick a hypo that generate the less amount of error. this is the learning algorithm

error of 0

|H| linear |D|

Theorem 2. For $\epsilon > 0$, the probability that the empirical error of the chosen hypothesis is off (the actual error) by more than ϵ is bounded from above by a decreasing function of $|D|$ and an increasing function of $|\mathcal{H}|$:

actual error and empirical error, difference in upper bound, however now theres the |H| as a coefficient

$$\mathbb{P}[|E(h^*) - E_D(h^*)| > \epsilon] \leq 2|\mathcal{H}|e^{-2\epsilon^2|D|}$$

probalistic statement that we dont have too large H or it becomes hard to learn

special case of hoeffding equality

Proof. The event $|E(h^*) - E_D(h^*)| > \epsilon$ implies:

as you have more hypo, you have more |H| so upper bound goes up. way to bring it down is to have more data - so want good balance between having lot data to have more confidence in upper bound - bringing it closer

$$\mathbb{P}[|E(h_1) - E_D(h_1)| > \epsilon] \vee \mathbb{P}[|E(h_2) - E_D(h_2)| > \epsilon] \vee \dots \vee \mathbb{P}[|E(h_{|\mathcal{H}|}) - E_D(h_{|\mathcal{H}|})| > \epsilon]$$

if we have lots of H, then want lots of data to help

that is

$$\mathbb{P}[|E(h^*) - E_D(h^*)| > \epsilon] \subseteq \bigcup_{h \in \mathcal{H}} \mathbb{P}[|E(h) - E_D(h)| > \epsilon]$$

Thus,

$$\mathbb{P}[|E(h^*) - E_D(h^*)| > \epsilon] \leq \mathbb{P}\left[\bigcup_{h \in \mathcal{H}} \mathbb{P}[|E(h) - E_D(h)| > \epsilon]\right]$$

By Boole's inequality:

$$\leq \sum_{h \in \mathcal{H}} \mathbb{P}[|E(h) - E_D(h)| > \epsilon]$$

where by applying Hoeffding's inequality to each summand, we get

$$\leq 2|\mathcal{H}|e^{-2\epsilon^2|D|}$$

□